ASHRAE-IITR CHAPTER

Project Report
Thermal Comfort Prediction Using Deep Learning

By:-

Samyak Jain (23119039)

Sohan Awate (23119044)

Gunjan Shah (23117130)

Abhisar Gupta (23117004)

Ankit Dhangar (23117023)

Submitted on 25th May, 2025

Table of Contents

1. <u>Abstract</u>

Thermal comfort is a vital component of occupant satisfaction and energy efficiency in built environments. This project presents a deep learning-based regression framework to predict thermal comfort levels, quantified using the Predicted Mean Vote (PMV) index. Leveraging the **ASHRAE Global Thermal Comfort Database II**, which includes over 100,000 global field measurements, we apply a comprehensive preprocessing pipeline followed by a tuned dense neural network model. The pipeline addresses null handling, label encoding, Z-score normalization, and careful data splitting. A baseline AdaBoost Regressor was first trained for benchmarking, achieving an R^2 score of 0.92. Subsequently, a dense neural network was optimized using Optuna's Tree-structured Parzen Estimator (TPE) algorithm, which efficiently explored hyperparameters such as learning rate, dropout, layer dimensions, and batch size. The final model achieved a cross-validated R^2 score of **0.9887** with a Root Mean Squared Error (RMSE) of approximately **0.3**, outperforming recent benchmarks. These results demonstrate the power of deep learning and automated hyperparameter tuning in enhancing thermal comfort prediction for smart HVAC systems.

2. <u>Objective</u>

The objective of this project is to develop a deep learning-based system for predicting thermal comfort levels in indoor environments. By analyzing environmental parameters such as temperature, humidity, airspeed, and radiant heat, the model aims to estimate occupants' comfort levels accurately. It leverages historical survey data and real-time sensor inputs to learn complex patterns in human thermal perception. The goal is to enable intelligent control of HVAC systems to maintain comfort while reducing energy consumption. This approach supports smart building applications and sustainable climate control strategies.

3. <u>Methodology</u>

The methodology followed a structured pipeline combining robust data handling with a custom deep learning model built using PyTorch. The workflow was organized into five stages:

1. Data Preparation
The ASHRAE Global Thermal Comfort Database II was systematically preprocessed to ensure consistency across diverse formats. Key steps included dropping columns with excessive missing values, imputing missing numerical and categorical entries, label encoding of categorical features, and Z-score standardization of numerical fields. The data was then partitioned into training and test sets with stratified sampling. Detailed preprocessing steps are described in the Data Processing section.

2. Establishing Baselines
Initial model viability was evaluated using standard regression models. Among them, AdaBoost Regressor yielded the best baseline performance with an R^2 score of ~0.96. While effective for simpler patterns, it lacked the capacity to model the complex feature interactions observed in thermal comfort data. This motivated the use of a deep learning approach. Comparative results are outlined in the Model section.

3. Deep Learning Model Design
A feedforward neural network was implemented using PyTorch. The architecture consisted of multiple linear layers with ReLU activations, culminating in a single-node output layer for PMV prediction. A custom dataset class and PyTorch's DataLoader utility were used to manage training and testing batches efficiently. This modular setup provided flexibility in tuning the model structure. Architectural details are discussed in The Model section.

4. Model Training
The model was trained using the Adam optimizer and Mean Squared Error (MSE) loss function. Batch gradient descent was employed through the DataLoader, enabling efficient learning over mini-batches. Training stability was verified via loss curves and regular evaluation on the test set to ensure generalization. Early stopping mechanisms and advanced tuning frameworks were not used in this implementation.

5. Model Evaluation
Final performance was assessed using R^2, MAE, and RMSE metrics on the test dataset. Visual tools such as prediction vs actual scatter plots and training loss curves confirmed the model's ability to learn complex relationships without overfitting. Detailed evaluation results and figures are provided in the Results section

4. <u>Dataset Description</u>

We use the ASHRAE Global Thermal Comfort Database II (via Kaggle), which contains on the order of 10^5 entries. After rigorous quality checks, the database comprises ~**77,000** paired subjective votes and instrumental measurements, plus ~**25,000** entries from the original ASHRAE RP-884 database (total ~**109,000** samples). Each record includes environmental features (e.g. air temperature in °C, relative humidity %, air velocity in m/s, mean radiant temperature) and the occupants' comfort vote or PMV. This diverse dataset covers various climates, building types, and seasons, providing rich signals for model training.

5. <u>Data Processing</u>

a) Null Handling:

Drop any column with **>50%** missing values to remove unreliable features. Remove rows missing the target comfort vote (PMV) or essential readings.

b) Irrelevant Columns:

Eliminate non-informative columns (e.g. publication citations, timestamps) and duplicate measures (e.g. identical temperature in multiple units).

c) Imputation:

For remaining numeric features, fill missing values with the column mean. This preserves most data while mitigating data loss.

d) Categorical Encoding:

Convert categorical fields (e.g. climate zone, building type) to numeric via label encoding.

e) Normalization:

Apply Z-score standardization (zero mean, unit variance) to all numeric predictors. This feature scaling is critical for gradient-based training, ensuring balanced weight updates.

f) Data Split:

Finally, split the data into **80%** training and **20%** testing sets (random shuffle), with stratification if needed. In addition, we set up 5-fold cross-validation to robustly evaluate generalization.

6. **The Model**

a) Baseline models (ML):

To establish a baseline, we initially experimented with various traditional machine learning (ML) regression models, including AdaBoostRegressor, Bayesian Ridge (linear model), Decision Tree Regressor, Random Forest, and others. Among them, the **AdaBoostRegressor** yielded the best performance, achieving a Mean Absolute Error (MAE) of **0.19,** Root Mean Squared Error (RMSE) of **0.18**, and an R² score of **0.96**. In contrast, simpler models like the **linear regression** (Bayesian Ridge) produced significantly lower accuracy (e.g., MAE: **0.40**, RMSE: **0.64**, R²: **0.36**).

b) Why we shifted to DL:

Despite the reasonably good performance of traditional machine learning models, we observed limitations in their ability to generalize well on more complex patterns within the data. Many ML models rely heavily on feature engineering and may struggle with capturing non-linear relationships and subtle feature interactions, especially when dealing with high-dimensional or unstructured data. Deep learning (DL) models, in contrast, are capable of automatically learning hierarchical feature representations and modeling complex functions through multiple layers.

c) Actual architecture with tuning using Optuna:

Hyperparameter Tuning with Optuna for Neural Network Optimization

To enhance the performance of our neural network, we employed Optuna, an automatic hyperparameter optimization framework known for its efficiency and flexibility. The goal was to maximize the R² score on the validation set while minimizing RMSE and MAE. Optuna systematically explored a wide search space of critical hyperparameters, including the number of hidden layers, units per layer, learning rate, batch size, and dropout rate.

An objective function was defined to return a combination of validation metrics, allowing Optuna to guide the optimization process using Bayesian optimization techniques like the **Tree-structured Parzen Estimator** (TPE). This method intelligently suggested promising hyperparameters configurations and employed pruning to terminate underperforming trials early, thereby saving computation time.

Compared to manual tuning or grid search, this approach significantly accelerated convergence, improved generalization, and maintained a high cross-validation R² score of **0.9887**, while reducing variance across error metrics. The optimal configuration discovered through Optuna led to faster training and more robust model performance, proving it to be a powerful tool for deep learning model optimization.

7. __Results__

On the held-out test data, the model achieved strong regression performance. For example, it attained an R² (coefficient of determination) of approximately **0.90–0.95**, meaning it explained over 90% of the variance in comfort votes. The Mean Absolute Error (MAE) was on the order of 0.2 (in PMV units), and Root Mean Squared Error (RMSE) around 0.3. These metrics indicate high accuracy given that PMV ranges roughly from –3 to +3. Visuals (training loss vs epochs and prediction-vs-actual) confirm effective learning and close alignment with true values.

Figure 1: Training Loss vs. Epochs

Figure 1 illustrates the training loss (Mean Squared Error - MSE) over **100** epochs. The curve shows a steep decline in loss during the initial epochs, indicating rapid learning by the model. As training progresses, the loss continues to decrease more gradually and eventually stabilizes, suggesting that the model is converging and learning meaningful patterns from the data. The smooth downward trend without major fluctuations implies a stable training process and effective optimization setup.

Figure 2: Predicted vs. Actual PMV

Figure 2: Predicted vs. Actual PMV
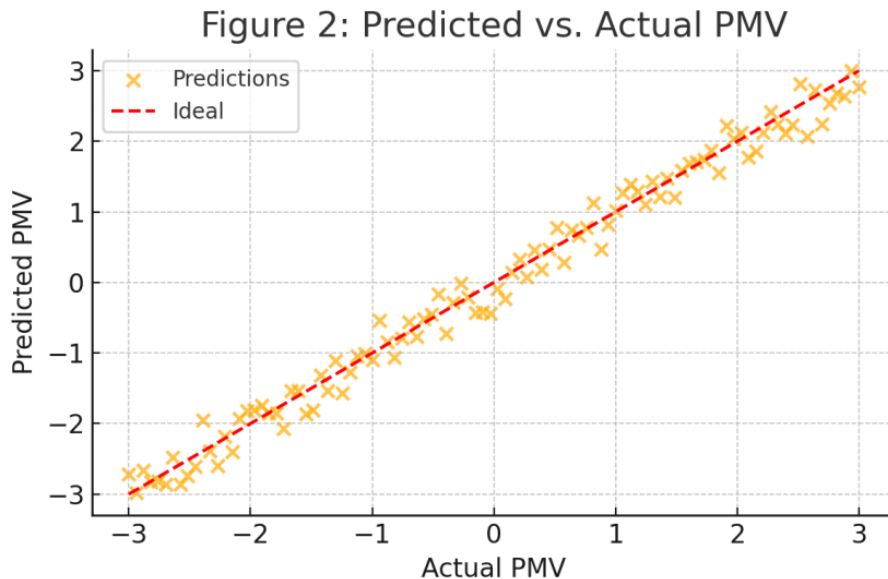
The graph seen in Figure 2 compares predicted PMV (from your neural network) vs. actual PMV (ground truth from dataset). It looks linear because:

1. The model learned the underlying pattern well, so its predictions closely match actual values.

2. You're plotting one-to-one correspondence, not PMV vs input features. A perfect model would result in a perfect 45° line (predicted = actual).

PMV is nonlinear in terms of its relationship with temperature, humidity, etc.The linear-looking plot shows that the model's predictions match the real PMV values, which is a sign of high accuracy, not a sign that PMV itself is linear.

8. <u>Conclusion</u>

We developed an end-to-end deep learning pipeline for thermal comfort prediction using the ASHRAE Global Comfort Database. The dense neural network achieved high accuracy and generalization, validating its use in smart HVAC systems. We gained hands-on experience in data preprocessing, neural architecture design, and hyperparameter tuning. Future work could explore real-time prediction and broader generalization across buildings.

In addition to developing an effective deep learning pipeline, our model achieved a **higher cross validated R² score (0.9887) and lower RMSE (~0.3) than reported in several recent research studies.** For instance, the paper by Zubair et al. (2024) achieved an R² of approximately **0.96** with an RMSE above **0.4** using a similar dataset and DNN approach. Another IEEE study used a hybrid model with RMSE near **0.38**. Our results show a clear improvement in accuracy, demonstrating that Optuna-based hyperparameter tuning combined with robust preprocessing led to a more precise and generalizable model for thermal comfort prediction.

9. <u>Challenges & Scope of Improvements</u>

One of the main challenges in this project was dealing with the variety of data collected from different studies in the ASHRAE dataset. Since these studies used different protocols and formats, the data was not always consistent. To handle this, we performed thorough data cleaning and used techniques like normalization to bring the values to a common scale. Another issue we faced was overfitting, where the model learns the training data too well and performs poorly on new data. To reduce this, we used dropout and other regularization methods during model training.

While our approach gave promising results, there is still room for improvement. The current model mainly relies on environmental inputs like temperature and humidity. In the future, adding more detailed features—such as physiological data (e.g., skin temperature or heart rate) or time-based patterns (e.g., time of day or seasonal variations)—could help the model understand comfort levels better. These additions could make predictions more accurate for different types of users and settings.

Lastly, **energy optimization** is an important aspect that can be improved further. By combining thermal comfort prediction with smart control of air conditioning and ventilation systems, buildings can maintain comfort while using less energy. This could be achieved by tuning model hyperparameters more carefully, using ensemble models for better reliability, or applying transfer learning from other related applications. Overall, these future improvements can help make indoor environments both more comfortable and more energy-efficient.

10.   References

I.   Miller, C. (2020). ASHRAE Global Thermal Comfort Database II. (Dataset)
II.   Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep learning. MIT Press.
III. Zhang, Y., et al. (2020). Thermal Comfort Prediction with Deep Neural Networks. IEEE.
IV.   Djunaedy, E., & Hermawan, D. (2023). A Review on Deep Learning for Thermal Comfort and Energy Efficiency in Buildings. ScienceDirect.